

CS234 - Reinforcement Learning

A Markov Reward Process (**MRP**) is a Markov chain + rewards (everything only depends on the state $R(s_t = s)$). The **total reward** (the **return**) is

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^{H-1} r_{t+H-1} = \sum_{k=0}^{H-1} \gamma^k r_{t+k}.$$

Expected return:

$$V(s) = \mathbb{E}[G_t | s_t = s] = R(s) + \gamma \sum_{s' \in S} P(s' | s) V(s').$$

Let a N -dim statespace, $V = (V(s_1), \dots, V(s_N)) \in \mathbb{R}^N$, $R, V \in \mathbb{R}^N$, and $P \in \mathbb{R}^{N \times N}$ with $P_{ij} = P(s_j | s_i)$. Then, we can write the matrix equation $V = R + \gamma P V$ and invert it $V = (I - \gamma P)^{-1} R$. Complexity of doing matrix inverse is $O(N^3)$. Otherwise, do, iteratively over k , for all $s \in S$,

$$V_k(s) = R(s) + \gamma \sum_{s' \in S} P(s' | s) V_{k-1}(s'), \quad (O(N^2) \text{ complexity})$$

A Markov Decision Process (**MDP**) is a 5-tuple (S, A, P, R, γ) where P is a transition model $P(s_{t+1} = s' | s_t = s, a_t = a)$ and R is reward function $R(s, a)$ and $\gamma \in [0, 1]$ is discount factor. **State-Value Function** (expected return under π from s):

$$V^\pi(s) = \mathbb{E}_\pi[G_t | s_t = s] = \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s].$$

Optimal policy: $\pi^* = \arg \max_\pi V^\pi(s)$.

State-Action Value Function (expected return starting from s , taking action a , then following π):

$$\begin{aligned} Q^\pi(s, a) &= \mathbb{E}_\pi[G_t | s_t = s, a_t = a] \\ &= \mathbb{E}_\pi[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a] \\ &= R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^\pi(s') \end{aligned}$$

Bellman Backup:

$$\begin{aligned} (BV)(s) &= \max_a (R(s, a) + \gamma \sum_{s'} P(s' | s, a) V(s')) \\ (B^\pi V)(s) &= \sum_a \pi(a | s) \left(R(s, a) + \gamma \sum_{s'} P(s' | s, a) V(s') \right) \end{aligned}$$

Policy Evaluation(π):

$$\begin{aligned} &V_0^\pi(s) \leftarrow 0 \text{ for all } s \\ &\textbf{while } \|V_{k+1}^\pi - V_k^\pi\| > \epsilon \textbf{ do} \\ &\quad \textbf{forall } s \in S \textbf{ do} \\ &\quad \quad V_{k+1}^\pi(s) = (B^\pi V_k^\pi)(s). \end{aligned}$$

Policy Improvement(V^π):

$$\begin{aligned} Q^\pi(s, a) &= R(s, a) + \gamma \sum_{s'} P(s' | s, a) V^\pi(s') \quad \forall s \in S, a \in A \\ \pi'(s) &= \arg \max_a Q^\pi(s, a) \quad \forall s \in S \end{aligned}$$

Policy Iteration:

$$\begin{aligned} &\pi_0(s) \leftarrow \text{whatever}(A) \text{ for all } s \\ &\textbf{while } \|\pi_{k+1} - \pi_k\| > \epsilon \textbf{ do} \\ &\quad V_k^\pi \leftarrow \text{PolicyEvaluation}(\pi_k) \\ &\quad \pi_{k+1}(s) \leftarrow \text{PolicyImprovement}(V_k^\pi) \end{aligned}$$

Value Iteration:

$$\begin{aligned} &V_0(s) \leftarrow 0 \text{ for all } s \\ &\textbf{while } \|V_{k+1} - V_k\| > \epsilon \textbf{ do} \\ &\quad \textbf{forall } s \in S \textbf{ do} \\ &\quad \quad V_{k+1}(s) = (BV_k)(s) \end{aligned}$$

These algorithms assume Markov and require knowledge of the dynamics P . Instead, we can do (model-free) **Monte-Carlo** (requires fully observed finite episode). It is unbiased (assumes infinite-horizon problem (so that time-invariant policy) that is always eventually stopped (at the stopping time T)).

Monte-Carlo(π)

$$\begin{aligned} &N(s) = 0, G(s) = 0, V^\pi(s) = 0 \text{ for all } s \\ &\textbf{for } i = 0, \dots \textbf{ do} \\ &\quad \text{sample episode } (s_{i,1}, a_{i,2}, \dots, s_{i,T_i}) \text{ with } a_{i,t} \sim \pi(s_t) \\ &\quad G_{i,t} \leftarrow r_{i,t} + \gamma r_{i,t+1} + \dots + \gamma^{T_i-t} r_{i,T_i} \text{ for all } t \\ &\quad \textbf{for } t = 1, \dots, T_i \textbf{ do} \\ &\quad \quad s \leftarrow s_{i,t} \\ &\quad \quad \textbf{if } s \text{ visited for the first time in episode } i \textbf{ then} \\ &\quad \quad \quad N(s) = N(s) + 1 \\ &\quad \quad \quad G(s) = G(s) + G_{i,t} \\ &\quad \quad \quad V^\pi(s) = G(s)/N(s) \\ &\quad \quad \quad (V^\pi(s) \leftarrow V^\pi(s) + \alpha(G_{i,t} - V^\pi(s)) \text{ (Incremental)}) \\ &\quad \quad \quad (\Delta w = \alpha(G_{i,t} - \phi(s_{i,t})^\top w) \phi(s_{i,t}) \text{ (MC linear)}) \end{aligned}$$

The above is **first-visit** MC. Without the *if*, it is **every-visit** MC (which is biased but consistent and has better MSE).

Modification to get Q^π : use $N(s, a) = N(s, a) + 1$ and set $Q^\pi(s, a) = Q^\pi(s, a) + \alpha(G_{i,t} - Q^\pi(s, a))$ with $\alpha = \frac{1}{N(s, a)}$.

TD(0)(π) (instead of waiting for the full episode to finish)

$$\begin{aligned} &V^\pi(s) = 0 \text{ for all } s \\ &\textbf{for } i = 0, \dots \textbf{ do} \\ &\quad \text{sample tuple } (s_t, a_t, r_t, s_{t+1}) \text{ with } a_t \sim \pi(s_t) \\ &\quad V^\pi(s_t) \leftarrow V^\pi(s_t) + \alpha \left(\underbrace{(r_t + \gamma V^\pi(s_{t+1}))}_{\text{TD target (bootstrap)}} - V^\pi(s_t) \right) \\ &\quad (\Delta w = \alpha(r_t + \gamma \phi(s_{t+1})^\top w - \phi(s_t)^\top w) \phi(s_t) \text{ (TD(0) linear)}) \end{aligned}$$

Certainty-Equivalence (CE) (MLE estimate for P and R)

$$\begin{aligned} &\textbf{for } t = 0, \dots \textbf{ do} \\ &\quad \text{sample tuple } (s_t, a_t, r_t, s_{t+1}) \text{ with } a_t \sim \pi(s_t) \\ &\quad \hat{P}(s' | s, a) = \frac{1}{N(s, a)} \sum_{r=1}^t \mathbb{1}(s_r = s, a_r = a, s_{r+1} = s') \\ &\quad \hat{r}(s, a) = \frac{1}{N(s, a)} \sum_{\ell=1}^t \mathbb{1}(s_\ell = s, a_\ell = a) r_\ell \end{aligned}$$

SARSA and **Q-Learning** with ϵ -greedy policy

$$\begin{aligned} &a_0 \sim \pi(s_0) \text{ w.p. } 1 - \epsilon \text{ else } a_0 \leftarrow \text{random}(A) \\ &\text{Observe } (r_0, s_1) \\ &\textbf{for } t = 0, \dots \textbf{ do} \\ &\quad \text{Take } a_{t+1} \sim \pi(s_{t+1}), \text{ observe } (r_{t+1}, s_{t+2}) \\ &\quad Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)) \\ &\quad Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)) \\ &\quad \pi(s) \leftarrow \{\arg \max_a Q(s, a) \text{ w.p. } 1 - \epsilon \text{ else } \text{random}(A)\} \end{aligned}$$

SARSA (an *on-policy* algorithm) uses every element of the tuple $(s_t, a_t, r_t, s_{t+1}, a_{t+1})$. **Q-Learning** (*off-policy*) allows learning optimum from arbitrary policies (only used for

exploration) but suffers from the maximization bias. Thus, **Double Q-Learning** separates max action estimation with max value estimation: with probability $1/2$, alternate between Q_1 and Q_2 : for 1, set $a_{t+1}^* = \arg \max_a Q_1(s_{t+1}, a)$ and evaluate $Q_1(s_t, a_t) \leftarrow Q_1(s_t, a_t) + \alpha(r_t + \gamma Q_2(s_{t+1}, a_{t+1}^*) - Q_1(s_t, a_t))$ $\pi(s) \leftarrow \{\arg \max_a (Q_1 + Q_2)(s, a) \text{ w.p. } 1 - \epsilon \text{ else } \text{rand}(A)\}$

Value function approximation

S and A are assumed finite-dim. Define the loss $J(w) = \mathbb{E}_\pi[(V^\pi(s) - \hat{V}^\pi(s; w))^2]$. Then, supervised learning: $\Delta w \stackrel{\text{SGD}}{=} -\frac{1}{2} \alpha \left(2\mathbb{E}_\pi[V^\pi(s) - \hat{V}^\pi(s; w)] \nabla_w \hat{V}^\pi(s; w) \right)$

Linear: $\hat{V}(s; w) = \phi(s)^\top w \implies \nabla_w \hat{V}(s; w) = \phi(s)$.

MC π evaluation: dataset $\{(s_t, G_t)\}_t$, use $V^\pi(s_t) \approx G_t$

Converges to $\min \text{MSVE}_\mu(w) = \sum_s \mu(s) (V^\pi(s) - \hat{V}^\pi(s; w))^2$

TD π evaluation: dataset $\{(s_t, r_t + \gamma \hat{V}^\pi(s_{t+1}; w))\}_t$

Let $d(s)$ stationary distribution s.t. $d(s') = \sum_s \sum_a \pi(a | s) p(s' | s, a) d(s)$

Converges to $\min \text{MSVE}_d(w) \leq \frac{1}{1-\gamma} \min_w \sum_s d(s) (V^\pi(s) - \hat{V}^\pi(s; w))^2$

SARSA, Q-Learning, and MC control: do the same for Q :

$$\Delta w = \alpha(r + \gamma \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

$$\Delta w = \alpha(r + \gamma \max_{a'} \hat{Q}(s', a'; w) - \hat{Q}(s, a; w)) \nabla_w \hat{Q}(s, a; w)$$

$$\Delta w = \alpha(G_t - \hat{Q}(s_t, a_t; w)) \nabla_w \hat{Q}(s_t, a_t; w)$$

Instability: (1) function approx (2) bootstrapping (3) off-policy

Deep Q-Learning: (1) correlated samples (2) non-stationary targets are addressed via (a) **experience replay** (b) **fixed Q-targets**

$$\begin{aligned} &\textbf{for } t = 0, \dots \textbf{ do} \\ &\quad \text{Sample } a_t \text{ given } \epsilon\text{-greedy policy from } \hat{Q}(s_t, a; w) \\ &\quad \text{Observe } (r_t, s_{t+1}) \\ &\quad \text{Store } (s_t, a_t, r_t, s_{t+1}) \text{ in replay buffer } D \\ &\quad \text{Sample minibatch from } D \\ &\quad \textbf{for tuple in minibatch do} \\ &\quad \quad \Delta w = \alpha(r + \gamma \max_{a'} \hat{Q}(s', a'; w^-) - \hat{Q}(s, a; w)) \nabla \hat{Q}(s, a; w) \\ &\quad \quad \textbf{if sometimes then} \\ &\quad \quad \quad w^- \leftarrow w \end{aligned}$$

Double DQN even better

$$\Delta w = \alpha(r + \gamma \hat{Q}(s', \arg \max_{a'} \hat{Q}(s', a'; w); w^-) - \hat{Q}(s, a; w)) \nabla \hat{Q}(s, a; w)$$

Prioritized Experience Replay: select i th tuple such that

$$p_i = |r + \gamma \max_{a'} Q(s', a'; w^-) - Q(s, a; w)| \quad \forall i \text{th tuple in } D$$

then choose i th tuple $(s, a, r, s') \in D$ for Δw update with $\max p_i$.

Else, select with stochastic prioritization, where $P(i) = p_i^\beta / \left(\sum_j p_j^\beta \right)$.

Advantage Function (Dueling DQN) Features to represent V well may be different than those to compare $Q(s, a_1)$ vs $Q(s, a_2)$.

$$\text{Define } A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s). \text{ Then, to get } \hat{Q}:$$

$$\begin{aligned} \hat{Q}(s, a; w) &= \hat{V}(s; w) + \left(\hat{A}(s, a; w) - \max_{a'} \hat{A}(s, a'; w) \right) \\ \hat{Q}(s, a; w) &= \hat{V}(s; w) + \left(\hat{A}(s, a; w) - \frac{1}{A} \sum_{a'} \hat{A}(s, a'; w) \right) \end{aligned}$$

Summary:

• Model-based RL

- + *easy* to learn a model (supervised learning)
- + learns *all there is to know* from the data
 - uses compute & capacity on irrelevant details
 - computing π (planning) non-trivial / expensive

• Value-based RL

- + easy to get policy (e.g., $\pi(a|s) = \mathbb{1}(a = \arg \max Q(s, a))$)
- + close to true objective
- + fairly well-understood, good algorithms exist
 - still not the true objective (may focus capacity on irrelevant details, small value error can lead to larger policy error)

• Policy-based RL

- + true objective
- + easy to extend to high-dim or continuous A
- + can learn stochastic policies
- + policies can be simple compared to values / models
 - local minima
 - does not always generalise well

Value Based


- Learn values
- Implicit policy (e.g. ϵ -greedy)

Policy Based

- No values
- Learn policy

Actor-Critic

- Learn values
- Learn policy



Policy learning: directly learn $\pi_\theta(s, a)$ minimizing $V(\theta)$:

$$\begin{aligned} V(s_0, \theta) &= \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^T R(s_t, a_t) \right] = \mathbb{E}_{\pi_\theta} [Q(s_0, \cdot, \theta)] \\ &= \sum_{\tau=(s_0, a_0, r_0, \dots, r_{T-1}, s_T)} P(\tau; \theta) R(\tau). \end{aligned}$$

Policy Gradient Theorem: if π_θ differentiable,

$$\nabla_\theta V(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, \cdot) Q(s_0, \cdot, \theta)].$$

Its gradient can be approximated as

$$\begin{aligned} \nabla_\theta V(\theta) &= \nabla_\theta \sum_{\tau} P(\tau; \theta) R(\tau) = \sum_{\tau} R(\tau) \nabla_\theta P(\tau; \theta) \\ &= \sum_{\tau} R(\tau) P(\tau; \theta) \nabla_\theta \log P(\tau; \theta) \\ &\approx \frac{1}{M} \sum_{i=1}^M R(\tau^i) \nabla_\theta \log P(\tau^i; \theta) \\ &= \frac{1}{M} \sum_{i=1}^M R(\tau^i) \left(\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \right), \end{aligned}$$

where the last step follows from the fact that

$$P(\tau^i; \theta) = \mu(s_0) \prod_{t=0}^{T-1} \pi_\theta(a_t | s_t) P(s_{t+1} | s_t, a_t).$$

Another derivation: if we exploit the temporal structure, then reward at time t doesn't depend on future timesteps:

$$\nabla_\theta \mathbb{E}[r_t] = \mathbb{E} \left[r_t \sum_{\tau=0}^t \nabla_\theta \log \pi_\theta(a_\tau | s_\tau) \right].$$

Thus,

$$\begin{aligned} \nabla_\theta V(\theta) &= \mathbb{E}_\tau \left[\sum_{t'=0}^{T-1} r_{t'} \left(\sum_{t=0}^{t'} \nabla_\theta \log \pi_\theta(a_t | s_t) \right) \right] \\ &= \mathbb{E}_\tau \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \sum_{t'=t}^{T-1} r_{t'} \right] \\ &\approx \frac{1}{M} \sum_{i=1}^M \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t^i | s_t^i) G_t^i. \end{aligned}$$

Note that one can add a baseline to reduce variance:

$$\nabla_\theta V(\theta) = \mathbb{E}_\tau \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \left(\sum_{t'=t}^{T-1} r_{t'} - b(s_t) \right) \right].$$

Does not introduce bias. Near optimal choice is the expected return $b(s_t) \approx \mathbb{E}[r_t + r_{t+1} + \dots + r_{T-1}]$.

Softmax policy: $\pi_\theta(s, a) = e^{\phi(s, a)^\top \theta} / \left(\sum_a e^{\phi(s, a)^\top \theta} \right)$.

Score function: $\nabla_\theta \log \pi_\theta(s, a) = \phi(s, a) - \mathbb{E}_{\pi_\theta} [\phi(s, a)]$.

Gaussian policy: $\pi_\theta(a(s)) \sim \mathcal{N}(\mu_\theta(s), \sigma^2)$.

Score function: $\nabla_\theta \log \pi_\theta(s, a) = \frac{(a - \mu_\theta(s))}{\sigma^2} \nabla \mu_\theta(s)$

Reinforce: uses $\Delta \theta_t = \nabla_\theta \log \pi_\theta(s_t, a_t) G_t$

for each episode $(s_1, a_1, r_1, \dots, s_{T-1}, \dots, r_T) \sim \pi_\theta$ **do**
 for $t = 1, \dots, T-1$ **do**
 $\Delta \theta = \alpha \nabla_\theta \log \pi_\theta(a_t | s_t) G_t$

Baseline $b(s)$ to reduce variance:

$$\theta_{t+1} = \theta_t + \alpha(r_{t+1} - b(s_t)) \nabla_\theta \log \pi_\theta(a_t | s_t)$$

Actor-Critic: use $\hat{V}^\pi(s; w)$ (estimated) as baseline

1-step Actor-Critic:

for $t = 1, \dots, T-1$ **do**
 Sample $a_t \sim \pi_\theta(a | s_t)$, observe (r_t, s_{t+1})
 $\delta_t = r_t + \gamma \hat{V}_w(s_{t+1}) - \hat{V}_w(s_t)$ (1step TD error/advantage)
 $w \leftarrow w + \beta \delta_t \nabla_w \hat{V}_w(s_t)$ (TD(0))
 $\theta \leftarrow \theta + \alpha \delta_t \nabla_\theta \log \pi_\theta(a_t | s_t)$ (policy gradient)

Increasing stability (TRPO, PPO): add KL regularization to an older policy π_{old} : do policy gradient with $J(\theta) - \eta \text{KL}(\pi_{\text{old}}, \pi_\theta)$:

$$\text{KL}(\pi_{\text{old}} \| \pi_\theta) = \mathbb{E}_s \left[\int_A \pi_{\text{old}}(a | s) \log \frac{\pi_{\text{old}}(a | s)}{\pi_\theta(a | s)} da \right]$$